


Introduction to machine learning


Artificial intelligence

- Emulate human intelligence: how?



Artificial intelligence

- Encode human knowledge




(doesn't work)

3

Artificial intelligence

- Use large amounts of examples



2
(it works!)

4

Machine learning

- Extract information (*patterns*) from data
- Applications:
 - Predictions (e.g. energy demand, sales)
 - Recommendations (e.g. Netflix)
 - Anomaly detection (e.g. intrusions, virus mutations)
 - Classification (e.g. image recognition, medical diagnoses)
 - Ranking (e.g. Google search)
 - Decision making (e.g. AI, robotics, games, trading)

5

Machine learning

- When it is useful
 - Lack of human support (e.g. robot on Mars)
 - Difficulty in encoding human experience (e.g. speech recognition, vision, language)
 - Too dimensionally-large problems (web page ranking, personalized advertising)

Success of machine learning

- Datasets
 - Image datasets in the '90s: 100 images
 - Image datasets in 2012: 1,000,000 images
 - Key factors: storage, data collection (image search), validation (crowdsourcing)
- General-purpose GPU
 - 100x speed-up w.r.t. CPUs
 - Training times of complex models: 1-2 weeks (on GPUs!)

7

Dataset

- Collection of data
 - Observation/samples/features: $\{x_i\}$
 - Potentially noisy, correlated, redundant
 - Label/target/output: $\{y_i\}$
- Examples:
 - Images + type of depicted content
 - Medical exams + positive/negative diagnosis
 - Past stock market trend + future stock market trend
 - Points (x_1, x_2, \dots, x_n) of a function and values $y = f(x_1, x_2, \dots, x_n)$

4

8

Model

- Objective: find the law that associates x to y
- **Model**: family of functions, defined by a set of **parameters**
- Algorithm design:
 - Choose the model according to our knowledge/expectation of the dataset
 - Choose the model's parameters: **training**

9

Non-parametric models

- E.g. nearest-neighbor
 - Dataset: $\{x_i, y_i\}$
 - Unknown input z
 - Associate to z the value y_j such that:

$$|z - x_j| = \min_{x_i} |z - x_i|$$

Learning modalities

- Supervised learning
- Unsupervised learning
- Reinforcement learning

11

Learning modalities

- Supervised learning:
 - Model is trained from $\{x_i, y_i\}$ pairs
 - Pros:
 - Most common, most studied, simplest
 - Training quality monitoring
 - Cons:
 - «Overfit» risk: the model learns the dataset **too well**
 - Dataset *annotation* (collecting $\{y_i\}$)

Learning modalities

- Unsupervised learning:
 - Model is trained from $\{x_i\}$
 - Pros:
 - Uncovers hidden dynamics (unknown even to experts)
 - Annotation is not necessary
 - Cons:
 - We may not know the set of possible values for y_i
 - Current techniques not satisfactory
 - E.g. *clustering, anomaly detection*

13

Learning modalities

- Reinforcement learning:
 - State: x_t , action: a_t
 - *Reward*: r_t , new state: x_{t+1}
 - No dataset, but an *environment simulator*
 - Pros:
 - Similar to human learning
 - Annotation is not necessary
 - Cons:
 - Hard to establish link between past actions and rewards
 - Hard to define reward
 - E.g. games

7

14

Supervised learning: classification


- $y_i \in \{C_1, C_2, \dots, C_k\}$
- *Categorical/discrete* labels
- **Binary** classification: $k = 2$
 - E.g. positivity/negativity of medical diagnosis
 - E.g. identify static/moving pixels in video
- **Multiclass** classification: $k > 2$
 - E.g. identify an object in an image

15

Supervised learning: regression

- $y_i \in \mathbb{R}$
- *Continuous* labels
- Examples:
 - Predict stock market prices
 - Predict user age on YouTube
 - Predict home energy consumption from temperature, sensors, weather, etc.

Classification vs regression



- Classification **separates**
- Regression **approximates**

17

Training

- How to choose model parameters?
- **Cost/loss/error/energy** function
 - Evaluates how good a set of parameter values is
 - Lower value → better parameters
- Analytical approach:
 - Compute θ such that $\frac{\partial L}{\partial \theta} = 0$
 - Usually intractable
- Iterative approach:
 - Initialize parameters θ_0
 - Repeat until convergence:
 - Evaluate loss $L(\theta_i)$
 - Re-compute $\theta_{i+1} \leftarrow \text{update}(\theta_i)$
- A «swipe» on the whole dataset is called **epoch**

18

Dataset splits

- **Training set:** fraction of the dataset used for training
 - The training algorithm uses these data only
 - Model/parameter choices must be conditioned on these data only
- **Test set:** fraction of the dataset used for verifying the model
 - The training algorithm and the designer **MUST NOT** see these data!
 - Evaluates how good the trained model in reality: **generalizability**

19

Dataset splits

- Why can't we trust the training set?
- **Overfit:** over-adaptation of the model on the training set, causing loss of generalizability
 - Reasons: model too complex, dataset too small
- **Unbalance**
 - Example: your dataset has 90 dog images and 10 cat images
 - Your model, for any image, predicts a dog
 - Accuracy of your model: 90% (?)

10

20

How to detect overfit

- The training algorithm can only use the training set
- The test set can be used at the end of training
- The model may have overfitted, but we only find out after evaluating on the test set
 - And we can do nothing about it, because we can't use the test set to improve our model
 - **Problem:** how to detect overfit?

21

How to detect overfit

- **Solution:** divide the original training set into a smaller training set and a **validation set**
- The new training set is used to update parameters
- The validation set is used during training to detect overfit
 - It's fine, it was part of the original training set
- The test set is only used at the end of training

Where we'll go from here

- Linear regression and classification
- Neural networks
 - Mostly for classification
- Convolutional neural networks
 - Mainly, 2D data processing (e.g. images)
- Recurrent neural networks
 - Sequential data processing (e.g. text, audio)

23

What you need to know

- Elements of linear algebra: scalar product, matrix multiplications
- Elements of calculus: derivatives and gradients
- Python